


1. The pdf of a $\Gamma(2, 1)$ random variable is $f(z) = z \exp(-z)$, $z > 0$, and the pmf of a Poisson random variable X is $f_X(x) = \lambda^x e^{-\lambda} / x!$, $\lambda > 0$, $x = 0, 1, \dots$. Assuming that X_1, X_2, \dots, X_n is an i.i.d Poisson sample and that λ has a $\Gamma(2, 1)$ prior distribution, find the MAP estimate of λ and prove that what you find is actually a value that maximizes the posterior.
2. DHS Chapter 3, Problem 17.
Hints: Note that $x_i \in \{0, 1\}$ and the problem walks you through constructing MAP estimates. For part (b), the uniform distribution for $p(\theta)$ extends over $0 \leq \theta \leq 1$; also, use Bayes theorem.
3. DHS Chapter 4, Problem 2
4. We perform Parzen Window density estimation, using trapezoidal window functions given (in unnormalized form) in the figure below:



Choose $h = 1$. Assume that we have the following data

$$\mathcal{D}_{\omega_1} = \{0, 2, 5\}$$

$$\mathcal{D}_{\omega_2} = \{4, 7\}$$

- (a) Sketch or plot the Parzen window estimates of the pdfs $p(x|\omega_1)$ and $p(x|\omega_2)$. Please label pertinent values on both axes.
 - (b) Estimate the prior probabilities based on frequency of occurrence of the prototypes in each class.
 - (c) Use the estimates you have developed in above to find the decision boundaries and regions for a Bayes minimum-error classifier based on Parzen windows. Only the part of feature space where at least one density is nonzero need to be classified.
5. We perform k-nearest neighbor density estimation, using $k = 2$. Assume that you are given the following training set for a 2-class problem with one feature:

$$\mathcal{D}_{\omega_1} = \{2, 5\}$$

$$\mathcal{D}_{\omega_2} = \{4, 7\}$$

- (a) Sketch or plot the k-nearest neighbors estimates of the pdfs $p(x|\omega_1)$. Please label pertinent values on both axes. Also give the density estimates algebraically, for each region in feature space.
 - (b) Estimate the prior probabilities based on frequency of occurrence of the prototypes in each class.
 - (c) Use the estimates you have developed in above to find the decision boundaries and regions for a Bayes minimum-error classifier based on k-nearest neighbors.
 - (d) Derive a classifier based on using KNN as a discriminative technique that estimates $p(\omega_i|x)$ directly using nearest neighbors, and compare it to the classifier you obtained in 5c.
6. For the following classification problem, design a single-layer perceptron that has zero error on training set, by using the multiclass Perceptron update rule.

$$\mathcal{D}_{\omega_1} = \left\{ \mathbf{x}_1 = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \right\}$$

$$\mathcal{D}_{\omega_2} = \left\{ \mathbf{x}_2 = \begin{bmatrix} 1 \\ -1 \end{bmatrix} \right\}$$

$$\mathcal{D}_{\omega_3} = \left\{ \mathbf{x}_3 = \begin{bmatrix} -1 \\ -1 \end{bmatrix} \right\}$$

Use one-hot coding for classes, for example, ω_3 should be represented using the following vector

$$\mathbf{y}_3 = \begin{bmatrix} -1 \\ -1 \\ 1 \end{bmatrix}$$

- (a) Start with $\mathbf{W}(0) = \mathbf{0}_{2 \times 3}$ and choose $\eta(i) = 0.5$ in $\mathbf{W}(i+1) = \mathbf{W}(i) + \eta(i)\mathbf{x}(i)\mathbf{e}^T$. Do not use the augmented space and assume that the biases are always zero (no update for biases). Show all steps until your algorithm converges. It is alright if you use a computer or calculator to perform the matrix calculations, but you should write down all the steps, and should not write a computer program to yield the final results.
- (b) Now redo the previous procedure, but this time deal with each of the columns of \mathbf{W} as one perceptron, i.e. update each column (the weight associated with a

linear discriminant) separately, for example the first iteration becomes:

$$\mathbf{W}(0) = [\mathbf{w}_1(0) \ \mathbf{w}_2(0) \ \mathbf{w}_3(0)]$$

$$\mathbf{w}_1(1) = \mathbf{w}_1(0) + \eta e_1 \mathbf{x}(1)$$

$$\mathbf{w}_2(1) = \mathbf{w}_2(0) + \eta e_2 \mathbf{x}(1)$$

$$\mathbf{w}_3(1) = \mathbf{w}_3(0) + \eta e_3 \mathbf{x}(1)$$

where $e_i = y_{i1} - \text{sign}(\mathbf{w}_i(0)^T \mathbf{x}(1))$ is the difference between the i^{th} element of \mathbf{y}_1 (the target vector for \mathbf{x}_1) and the output of the i^{th} neuron/linear discriminant $\mathbf{w}_i(1)$.

Perform two epochs only. This is essentially to make you observe that a multicategory Perceptron algorithm is based on multiple binary problems.